

Evaluating Employee Communication Behaviors Using Large Language Models

Zhanchen Kang and Yu-Liang Lin

School of Information Science

University of Illinois Urbana-Champaign

Champaign IL, USA

zk11@illinois.edu yl198@illinois.edu

Abstract

Large language models (LLMs) are increasingly deployed in customer-facing service contexts, yet their sociolinguistic behavior—specifically how well they calibrate empathy to a customer’s emotional state—remains largely unevaluated. This report describes a research project investigating whether frontier LLMs systematically over- or under-empathize relative to real human support agents, and whether such miscalibration varies across industry domains. We use the Customer Support on Twitter dataset (Axelbrooke, 2017) as a naturalistic behavioral benchmark, extracting 600 complaint–response pairs stratified across two domains (airlines, telecom) and three frustration levels. We prompt three frontier models (GPT-4o, Claude 3.5 Sonnet, Gemini 1.5 Flash) with each complaint and compare their outputs against human agent responses using a two-dimensional Empathy Calibration Score (ECS). This checkpoint report details the motivation, related work, methodology, experiment design, and a three-week execution plan.

1 Introduction

This study addresses two primary questions:

1. Do frontier LLMs systematically over-empathize or under-empathize relative to trained human customer support agents, as measured by a calibration delta against a naturalistic behavioral benchmark?
2. Is this empathy miscalibration consistent across industry domains, or does it vary by context—specifically between airlines and telecom?

Together, these questions target a concrete deployment risk: enterprises selecting or prompting LLMs for customer service may unknowingly be introducing systematic tonal mismatch. A model that sounds empathetic in general benchmarks may

be tonally inappropriate in practice—either performing hollow over-apology, or presenting as cold and procedural when a customer is genuinely distressed.

2 Related Work

2.1 Sycophancy and Over-Helpfulness in LLMs

A well-documented failure mode of RLHF-trained models is sycophancy—the tendency to agree with or flatter users at the expense of accuracy or appropriate restraint (Malmqvist, 2024; Dahlgren Lindström et al., 2025). Kim and Khashabi (2025) found that LLMs alter responses to conform with user expectations across a range of tasks. Recent mechanistic work shows that sycophantic agreement and sycophantic praise are distinct, independently steerable behaviors encoded in separate linear directions in latent space (Vennemeyer et al., 2026)—suggesting miscalibration is not monolithic. Separately, Zewe (2026) found that over extended interactions, personalization features significantly increase the likelihood an LLM becomes sycophantic, noting that “evaluation methods are lagging behind” how people actually use these models. In a customer service context, sycophancy manifests not as factual capitulation but as tonal over-correction: excessive apology, performed urgency, and mismatched warmth.

2.2 Empathy Evaluation in NLP

The NLP empathy literature is growing but domain-narrow. A systematic review of 12 studies on LLMs and empathy found that 7 focused exclusively on medical contexts, with all studies evaluating ChatGPT-3.5 (Sorin et al., 2024). Kumar et al. (2026) found that LLMs nearly match expert reliability when judging empathic communication, but their analysis was limited to personal-problem support contexts rather than transactional service dialogues. Existing metrics—BLEU, ROUGE,

BERTScore—are inadequate for empathy evaluation, measuring lexical overlap rather than emotional alignment (Iyer et al., 2026). A recent metric framework by Hong et al. (2025) proposes combining sentiment alignment and semantic relevance into a composite empathy score, but assumes that empathy should always trend positive—inappropriate for contexts where matching a customer’s matter-of-fact tone is the correct response.

2.3 Customer Support NLP

The most closely related applied work is Zhu et al. (2025), which introduces a customer support conversation dataset grounded in structured support strategies and evaluates multiple LLMs on their ability to generate empathetic, coherent responses. Notably, the authors found that GPT-4o produces shorter, less emotionally rich dialogues than DeepSeek-R1. However, CSConv uses expert-rewritten conversations as the gold standard rather than raw human agent behavior, idealizing away the naturalistic calibration signal that makes our benchmark distinctive. No prior work, to our knowledge, uses real, unmodified paired complaint–response data as a behavioral ground truth for evaluating LLM tonal calibration.

2.4 The Gap

Across these three bodies of work, a consistent blind spot emerges. Sycophancy research measures opinion-conformity, not tonal appropriateness. Empathy research evaluates support-context warmth, not calibration to frustration level. Customer service NLP evaluates response quality against idealized rewrites, not observed human behavior. Our study sits at the intersection of all three and uses a unique data asset—naturalistic, multi-domain, paired conversational data—to fill it.

3 Methodology

The study follows a three-phase empirical design: (1) benchmark construction from real conversational data, (2) LLM response generation under a neutral prompt, and (3) calibration scoring and statistical comparison. No model training is required.

3.1 Benchmark Construction

We parse the Twitter Customer Support dataset (Axelbrooke, 2017) to extract single-turn complaint–response pairs, where a customer tweet is followed

directly by a verified brand account reply. We compute a Frustration Intensity Score (FIS) for each customer tweet using a composite of VADER sentiment, all-caps ratio, punctuation density, and presence of urgency vocabulary (“never again”, “unacceptable”, “immediately”). FIS is binned into three levels: low, moderate, and high. We retain two industry domains—airlines and telecom—and apply stratified sampling to produce a benchmark of 600 instances (100 per cell: 2 domains $\times$ 3 FIS levels).

3.2 LLM Response Generation

Each of the 600 customer complaint tweets is passed to three frontier models under an identical, minimal system prompt:

“You are a customer support agent. Respond to the following customer message in one to three sentences.”

No persona, no tone instruction, and no domain context is supplied. This deliberate minimalism ensures we measure the model’s default social calibration rather than a prompted one. The three models evaluated are GPT-4o, Claude 3.5 Sonnet, and Gemini 1.5 Flash. The human agent’s original response from the dataset serves as the behavioral reference (ground truth).

3.3 Empathy Calibration Score (ECS)

The ECS is a two-dimensional rubric designed to measure not absolute empathy, but calibration—how well a response’s empathetic register matches what the customer’s expressed frustration level warrants. Table 1 describes each dimension.

Dimension	Measurement	Weight at high FIS
Apology density	Normalized count of apology/regret terms (regex). Fully automated.	Higher weight
Tone-frustration fit	LLM-as-judge (GPT-4o) scores 1–5: does the response tone match the emotional register of the complaint?	Equal weight

Table 1: ECS dimensions, measurement approach, and weighting at high frustration levels.

Scores for each dimension are computed for both the LLM response and the human agent response. The calibration delta— $ECS(LLM) - ECS(human)$ —is the primary outcome variable. A positive mean delta indicates the model is systematically over-empathetic relative to trained human agents; a negative delta indicates under-empathizing. A distribution centered at zero indicates calibration.

3.4 Validation

Before scaling, the LLM-as-judge rubric is validated on a 30-instance pilot, comparing GPT-4o judge scores against two human raters. Acceptable agreement is defined as Cohen’s $\kappa \geq 0.60$. If the pilot fails this threshold, the rubric is revised and re-piloted before the full run proceeds.

3.5 Statistical Analysis

Two tests map to the two research questions:

1. **Systematic miscalibration per model:** For each model, a one-sample t -test tests whether the mean calibration delta Δ differs significantly from zero. We report the mean Δ with a 95% confidence interval and Cohen’s d as an effect-size measure. A positive mean Δ with $p < 0.05$ constitutes evidence of systematic over-empathizing relative to trained human agents.
2. **Domain comparison:** A two-way ANOVA (model $\times$ domain) on the calibration delta tests whether miscalibration is domain-invariant or whether certain models are well-calibrated for airline complaints but miscalibrated for telecom, or vice versa.

4 Experiments

4.1 Datasets

The primary dataset is the Customer Support on Twitter corpus (Axelbrooke, 2017), a CSV of over 3 million tweets and replies from major consumer brands including airlines (AmericanAir, Jet-Blue, SouthwestAir, AlaskaAir), telecom providers (sprintcare, Ask_Spectrum, TMobileHelp), and others. Each row includes tweet ID, author ID, text, and a field for the tweet being replied to, enabling thread reconstruction.

We use this dataset in two ways:

- **Test benchmark:** 600 filtered, stratified complaint–response pairs (as described

above). These are never seen by any model during generation.

- **Judge calibration:** 30 additional pairs for the LLM-as-judge pilot validation, drawn outside the 600 main instances.

No additional external datasets are required. The human agent responses embedded in the dataset serve as the behavioral ground truth, which is the key design advantage of this corpus over synthetic alternatives.

For a detailed full conversation sample please check Table 2. This dataset is not ordered in how conversation is done. Thus, we applied a program to generate the fully ordered conversation shown in Appendix A.

4.2 Baselines

As for baselines, we evaluate three frontier models as candidate systems: GPT-4o, Claude 3.5 Sonnet, and Gemini 1.5 Flash. All three are evaluated under the same minimal neutral prompt with no persona or tone instruction. The human agent responses from the Twitter dataset serve as the behavioral ground truth baseline.

4.3 Evaluation Metrics

Model performance is evaluated using the two-dimensional ECS described in the Methodology section. The primary outcome is the calibration delta $\Delta = ECS(LLM) - ECS(human)$, computed per instance and aggregated per model and domain. We also report apology density independently as a fully model-agnostic signal.

4.4 Results

Table 3 shows a brief detail about composite, apology density and tone fit scores. The tested two models are both over-empathized, however, driven by different reason.

For GPT 4o, its major driver for over-empathized is Apology Density. However, for Gemini, its major driver for over-empathized is Tone Fit.

Figure 1 shows a calibration distribution by model under different frustration level. The 0 lines represent how human would react to customers under different frustration level.

Figure 2

4.5 Analysis

Figure 1 illustrates the distributions of calibration error ($\Delta = ECS_LLM - ECS_human$) across differ-

tweet_id	text	response_tweet_id
1	@115712 I understand. I ... further assist.	2
2	@sprintcare and how do you propose we do that	None
3	@sprintcare I have sent several private ... as usual	1
4	@115712 Please send us a Private ... your profile.	3
5	@sprintcare I did.	4
6	@115712 Can you please send ... your account?	5,7
8	@sprintcare is the worst customer service	9,6,10

Table 2: Samples of a full conversation in "Customer Support on Twitter"(Axelbrooke, 2017).

Model	Composite Delta	Δ Apology_density	Δ Tone_fit	Driver
GPT 4o	0.0973	0.1201	0.0694	Apology_density
Gemini	0.1736	0.1072	0.2494	Tone_fit

Table 3: Model Composite Δ , Δ Apology Density, Δ Tone Fit and Driver.

ent frustration levels for both GPT-4o and Gemini. A value of zero indicates perfect calibration, while positive values reflect systematic overestimation by the model.

For GPT-4o, the distribution under low frustration is sharply centered around zero, suggesting strong calibration when user frustration is minimal. However, as frustration increases, the distributions shift rightward and become more dispersed, indicating both increased overestimation and higher variance. This trend suggests that GPT-4o’s calibration degrades under more challenging emotional contexts.

In contrast, Gemini exhibits a consistent rightward shift across all frustration levels, with even low-frustration cases showing noticeable overestimation. The effect becomes more pronounced under moderate and high frustration, where the distributions are both further shifted and broader compared to GPT-4o. This indicates that Gemini not only overestimates more systematically but also demonstrates reduced stability.

Overall, both models are sensitive to user frustration, with higher frustration leading to greater positive calibration bias. However, GPT-4o remains better calibrated and more stable across conditions, whereas Gemini shows larger and more variable deviations from human judgments.

We also analyze model calibration across domains and examine whether domain-specific effects interact with model choice. Figure 2 summarizes the mean calibration error (Δ) for two models across the airline and technology domains, with 95% confidence intervals.

Across both domains, we observe a consistent

gap between the two models: gpt4o exhibits lower calibration error than gemini. In the airline domain, gpt4o achieves a mean Δ of approximately 0.05, compared to 0.12 for gemini. This pattern persists in the technology domain, where gpt4o reaches 0.14 while gemini increases to 0.23. These results indicate that gpt4o is systematically better calibrated, regardless of domain.

In addition, both models show higher calibration error in the technology domain than in the airline domain. This suggests that the technology domain poses greater challenges for reliable confidence estimation, potentially due to increased ambiguity or broader knowledge requirements.

5 Conclusions

This work is designed to test how LLMs perform under different frustration level and whether LLMs will be over or under empathized and what is its driver. We answered the following question in this paper.

1. RQ1 — Do frontier LLMs systematically over- or under-empathize relative to human agents?

Both models show statistically significant positive calibration deltas, meaning both over-empathize relative to human agents. Also, there’s different tendency in models and frustration levels.

2. RQ2 — Is this miscalibration consistent across industry domains (airline vs. technology)?




Figure 1: Calibration Delta Distributions by Model and Frustration Level.




Figure 2: Model and Domain Relationship.

The two-way ANOVA finds a significant main effect of domain, meaning both models over-empathize more in the technology sector than in the airline sector. A plausible interpretation: technology-sector human agents in the dataset tend to write more procedural, lower-warmth replies, so LLM defaults diverge more from that baseline than from the warmer airline-agent style.

6 Limitations

Due to the token and API shortage, we limited our models on only GPT-4o and Gemini. More tested model would provide a better result and a more detailed explanation on whether the over-empathize is a universal problem or is it only related to these two models.

Meanwhile, due to the limitation of customer service data, we cannot include all real-world conversation between a human agent and a customer. Thus, there might be a potential bias amount these analyzed data, which might cause the model to react over-empathized.

References

- Stuart Axelbrooke. 2017. [Customer support on twitter](#).
- Adam Dahlgren Lindström, Leila Methnani, Lea Krause, Petter Ericson, Íñigo Martínez de Rituerto de Troya, Dimitri Coelho Mollo, and Roel Dobbe. 2025. [Helpful, harmless, honest? sociotechnical limits of AI alignment and safety through reinforcement learning from human feedback](#). *Ethics and Information Technology*, 27(2):28. Epub 2025 Jun 4.
- Yuna Hong, Bonhwa Ku, and Hanseok Ko. 2025. [Performance evaluation metrics for empathetic llms](#). *Information*, 16(11).
- Laya Iyer, Kriti Aggarwal, Sanmi Koyejo, Gail Heyman, Desmond C. Ong, and Subhabrata Mukherjee. 2026. [Heart: A unified benchmark for assessing humans and llms in emotional support dialogue](#).
- Sungwon Kim and Daniel Khashabi. 2025. [Challenging the evaluator: Llm sycophancy under user rebuttal](#).
- Aakriti Kumar, Nalin Pongpeth, Diyi Yang, Erina Farrell, Bruce L. Lambert, and Matthew Groh. 2026. [When large language models are reliable for judging empathic communication](#). *Nature Machine Intelligence*, 8(2):173–185.

Lars Malmqvist. 2024. [Sycophancy in large language models: Causes and mitigations](#).

Vera Sorin, Dana Brin, Yiftach Barash, Eli Konen, Alexander Charney, Girish Nadkarni, and Eyal Klang. 2024. [Large language models and empathy: Systematic review](#). *J Med Internet Res*, 26:e52597.

Daniel Vennemeyer, Phan Anh Duong, Tiffany Zhan, and Tianyu Jiang. 2026. [Sycophancy is not one thing: Causal separation of sycophantic behaviors in LLMs](#).

Adam Zewe. 2026. [Personalization features can make LLMs more agreeable](#). MIT News.

Jie Zhu, Huaixia Dou, Junhui Li, Lifan Guo, Feng Chen, Chi Zhang, and Fang Kong. 2025. [Evaluating, synthesizing, and enhancing for customer support conversation](#).

A Full Conversation Sample

This section shows a sample of a full conversation.

Customer: @sprintcare I have sent several private messages and no one is responding as usual

Sprintcare: @115712 I understand. I would like to assist you. We would need to get you into a private secured link to further assist.

Customer: @sprintcare and how do you propose we do that

Sprintcare: @115712 Can you please send us a private message, so that I can gain further details about your account?

Customer: @sprintcare I did.

Sprintcare: @115712 Please send us a Private Message so that we can further assist you. Just click 'Message' at the top of your profile.

Customer: @sprintcare is the worst customer service